

# Visual Analytics— Dealing with Big Data: Dimensionality Reduction

Dr. Ab Mosca (they/them)

Slides based off slides courtesy of Jordan Crouser (<https://jcrouser.github.io/>)

# Reminder

- Be prepared for prototype testing in class on Thursday!

# Plan for Today

- Dimension Reduction Techniques
  - PCA
  - MDS
  - t-SNE
  - UMAP

What we've  
been (mostly)  
worried about

## Channels: Expressiveness Types and Effectiveness Ranks

### ➔ **Magnitude Channels: Ordered Attributes**

Position on common scale



Position on unaligned scale



Length (1D size)



Tilt/angle



Area (2D size)



Depth (3D position)



Color luminance



Color saturation



Curvature



Volume (3D size)



### ➔ **Identity Channels: Categorical Attributes**

Spatial region



Color hue



Motion



Shape



Most

Effectiveness

Least

Same

Same

What if  $\dim(\text{data}) \gg \# \text{channels}$ ?

Ideas?

## Dealing with Many Dimensions

- **Current situation:** our data live in  $p$ -dimensional space where  $p \gg \#$  visual channels
- Odds are not all dimensions are equally useful  
→ we can reduce the number of dimensions without losing too much valuable information

# Dimension reduction

## Approach #1: Feature Elimination (i.e. Subset Selection)

- Throw out less useful/useless dimensions

Pros? Cons?

- Pros:
  - Relatively simple
  - Preserve interpretability
- Cons:
  - Don't gain any information from features you drop

## Dimension reduction

### Approach #2: Feature Extraction

- Create new features (dimensions) that are combinations of the old ones

Pros? Cons?

- Pros:
  - Since new features are all combinations of old features, we are not totally dropping data
- Cons:
  - Less interpretable, especially to non-experts




## Feature Extraction


### Principal Component Analysis

- Project data into a smaller space composed of most important (informative) components


Projection



Projection




Projection



**Big idea:** *transform* the data into a new space

$$[X_1 \quad X_2] \mapsto [Z_1]$$

Projection



$$[X_1 \quad X_2 \quad X_3 \quad X_4 \quad X_5] \mapsto [Z_1 \quad Z_2]$$

## Linear projection

- New features are **linear combinations** of original data:

$$Z_j = \sum_i^m \theta_{ij} X_i$$


- We get them by multiplying the *data matrix* by a *projection matrix*

$$\begin{bmatrix} Z_1 & Z_2 \end{bmatrix} = \begin{bmatrix} X_1 & X_2 & X_3 & X_4 & X_5 \end{bmatrix} \begin{bmatrix} \varphi_{1,1} & \varphi_{1,2} \\ \varphi_{2,1} & \varphi_{2,2} \\ \varphi_{3,1} & \varphi_{3,2} \\ \varphi_{4,1} & \varphi_{4,2} \\ \varphi_{5,1} & \varphi_{5,2} \end{bmatrix}$$

## Why is projection helpful?

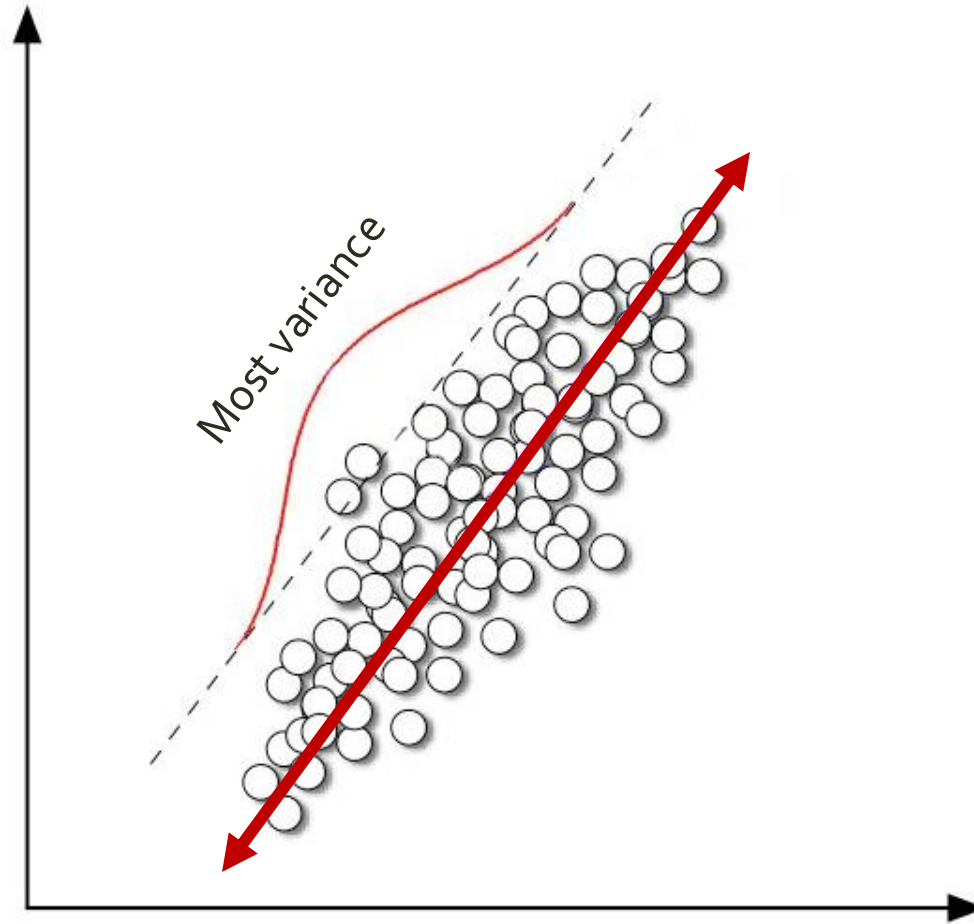
- Data can be rotated, scaled, and translated without changing the **underlying relationships**
- This means you're allowed to look at the data from whatever angle makes your life easier...
- Because new dimensions are combinations of old ones, we do not lose any data!

Flashback:  
why did we  
pick this line?






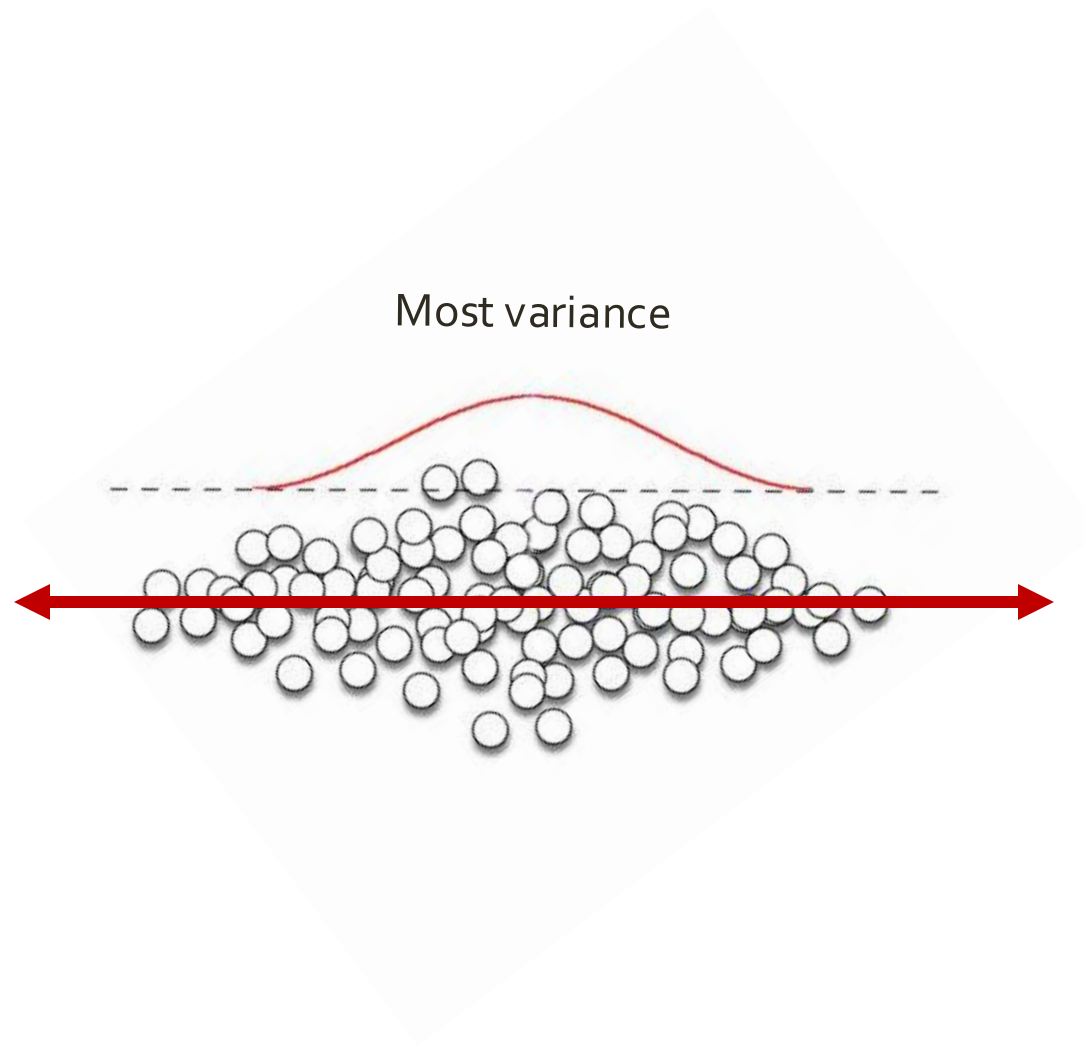
Explains the  
most **variance**  
in the data



Imagine this  
line as a new  
dimension...



“Principal  
component”



# Mathematically

- The **1<sup>st</sup> principal component** is the normalized\* linear combination of features:

$$Z_1 = \phi_{11}X_1 + \phi_{21}X_2 + \cdots + \phi_{p1}X_p$$

that has the largest variance

- $\phi_{11}, \dots, \phi_{p1}$ : the **loadings** of the 1<sup>st</sup> principal component


\* By **normalized** we mean:

$$\sum_{j=1}^p \phi_{j1}^2 = 1$$

## Using loadings to project

Multiply by loading vector to project (“smoosh”)  
each observation onto the line:

$$z_{i1} = \phi_{11}x_{i1} + \phi_{21}x_{i2} + \cdots + \phi_{p1}x_{ip}$$



These values are called the **scores**  
of the 1<sup>st</sup> principal component


## Additional principal components

- The **2<sup>nd</sup> principal component** is the normalized linear combination of the features

$$Z_2 = \phi_{12}X_1 + \phi_{22}X_2 + \cdots + \phi_{p2}X_p$$

that has maximal variance out of all linear combinations that are **uncorrelated** with  $Z_1$

Principal components are orthogonal



# Generating additional principal components

- We can think of this recursively
- To find the  $M^{th}$  principal component . . .
  - Find the first  $(M - 1)$  principal components
  - Subtract the projection into that space
  - Maximize the variance in the remaining *complementary* space



# Principal Component Analysis

- Conduct analysis using  $n$  principal components instead of the original data

How do we choose  $n$ ?

# Principal Component Analysis


- Conduct analysis using  $n$  principal components instead of the original data

Choosing  $n$ :

- **Option 1:** Arbitrarily choose
  - Ex. Will your audience only understand 2-d data?
- **Option 2:** Choose a proportion of data variance that must be captured by your principal components. Keep adding principal components until the threshold is hit.
- **Option 3:** Plot cumulative proportion of data variance captured by your principle components. Look for “elbow” where reduction in variance drops off.

# Principal Component Analysis

- **Option 3:** Plot cumulative proportion of data variance captured by your principle components. Look for “elbow” where reduction in variance drops off.



# PCA Exploration

- Take 10 minutes to explore using this applet:

<https://setosa.io/ev/principal-component-analysis/>


# Example

- Digits Dataset

Each datapoint is a 8x8 image of a digit.

Classes	10
Samples per class	~180
Samples total	1797
Dimensionality	64
Features	integers 0-16

Example: 1<sup>st</sup>  
two PCs from  
PCA



Potential issues  
with PCA?

## New idea

- Preserve high-dimensional pairwise distance in projection
  - “similar” stuff stays close together
  - “different” stuff can move apart




# MDS (Multidimensional Scaling)

1. Choose a good distance metric
2. Compute a pairwise distance matrix
3. Find a 2D embedding that preserves those distances- (the distance matrix kind of acts like a stress tensor, so you can think of MDS kind of like a “force directed” layout but without visible edges)



Fun (but also sad) fact:

- PCA is just a special case of MDS (if we use Euclidean distance and choose the 1st two components)

Example: 1<sup>st</sup>  
two PCs from  
MDS



# Example: 1<sup>st</sup> two PCs from PCA vs MDS




Potential issues  
with general  
MDS?

## t-SNE (t-Distributed Stochastic Neighbor Embedding)

- Same objective: preserve pairwise distances
- Different approach:
  - Similarity of points is determined by a probability distribution (t-distribution) and the conditional probability the point A will pick point B as it's "neighbor"
  - Density of points is taken into account via "perplexity"
  - Stochastic process (element of randomness)

Example: 1<sup>st</sup>  
two PCs from t-  
SNE



Potential issues  
with t-SNE?




# LDA (Linear Discriminant Analysis)





- Prioritizes class separability
- From a dataset of  $d$  independent features, extracts  $k$  new independent features that separate the classes the most
  - Note: you need to know classes to use LDA
- Algorithm is similar to PCA, but with the added constraints of minimizing inter-class spread and maximizing intra-class spread



Example: 1<sup>st</sup>  
two PCs from  
LDA



Potential issues  
with LDA?



# Takeaways

- There's no "one right answer"
- Each technique has benefits and drawbacks:
  - MDS (or PCA) is great if you eventually need to be able to relate the result back to the original dimensions
  - t-SNE does a great job preserving local similarity, but sometimes at the expense of global structure (and it can get really slow)
  - LDA is fast and relatively accurate